

# Introduction to XMPP



# What is XMPP?

- eXtensible Messaging and Presence Protocol
- Bi-directional streaming XML
- Core: IETF RFC [6120](#), [7590](#), [6121](#)
- Extensions: XMPP Standards Foundation (XSF)
  - Membership-based
  - Elected technical council
  - Unit of work: XMPP Extension Protocol (XEP)
  - Process: Experimental, Proposed, Draft, Final
- Goals:
  - Simple clients
  - Federate everything

# Related RFCs


| RFC      | Name      | Description  |
|----------|-----------|--|
| RFC 6120 | XMPP Core | XMPP core features<br>Updated by 7590 (TLS)  |
| RFC 6121 | XMPP IM   | XMPP Instant Messaging and Presence  |
| RFC 3922 | XMPP CPIM | Mapping XMPP to Common Presence and Instant Messaging (CPIM)                                   |
| RFC 3923 | XMPP E2E  | End-to-End Signing and Object Encryption for XMPP  |
| RFC 5122 | XMPP URI  | Internationalized Resource Identifiers (IRIs) and Uniform Resource Identifiers (URIs) for XMPP |
| RFC 4854 | XMPP URN  | Uniform Resource Name (URN) Namespace for XMPP   |

# What is XMPP ?

- The eXtensible Messaging and Presence Protocol (XMPP) is a TCP communications protocol based on XML that enables **near-real-time exchange of structured data** between two or more connected entities.
- Out-of-the-box features of XMPP include **presence information and contact list** maintenance.
- Due in part to its open nature and XML foundation, XMPP has been extended for use in **publish-subscribe systems**
  - **Perfect for IoT applications.**

<https://www.infoworld.com/article/2972143/internet-of-things/real-time-protocols-for-iot-apps.html>

# XMPP Architecture




Server<->Server: Port 5269

Client<->Server: Port 5222

# XMPP Architecture

- Addressing Scheme: `node@domain/resource`
  - JID = Jabber ID
  - Node: identity, e.g. **user name**
  - Domain: DNS **domain name**
  - Resource: **device identifier**
  - **`node@domain` identifies a person**
- Client talks to “local” server
  - Wherever the user account is hosted
  - Tied to directory if desired
  - Organizational policy enforced
- Servers talk to other servers
  - DNS lookup on domain portion of address
  - Dialback, MTLS for security
  - One connection for many conversations



# XML Refresher

- Element
- Attribute
- Namespace
- Language
- Text

Attribute

```
<geoloc xmlns='http://jabber.org/protocol/geoloc'
```

```
    xml:lang='en'
```

```
    id='14'>
```

Element

```
    <lat>38.9</lat>
```

```
    <lon>-77.1</lon>
```

```
    <locality>Arlington</locality>
```

```
    <region>VA</region>
```

```
</geoloc>
```

# XMPP Streams

- Client connects TCP socket to server

- Client sends stream start tag:

```
<stream:stream xmlns='jabber:client'  
               xmlns:stream='http://etherx.jabber.org/streams'  
               to='example.com'  
               version='1.0'>
```

- Server sends stream start tag back:

```
<stream:stream xmlns='jabber:client'  
               xmlns:stream='http://etherx.jabber.org/streams'  
               from='example.com'  
               id='someid'  
               version='1.0'>
```

- Each child element of stream: a “*stanza*”



# Stream features

- After stream start, server sends feature list:

```
<stream:features>  
  <starttls xmlns='urn:ietf:params:xml:ns:xmpp-tls' />  
  <mechanisms xmlns='urn:ietf:params:xml:ns:xmpp-sasl'>  
    <mechanism>DIGEST-MD5</mechanism>  
  </mechanisms>  
  <compression xmlns='http://jabber.org/features/compress'>  
    <method>zlib</method>  
  </compression>  
</stream:features>
```

- Client can negotiate any of these features

# XML Stream Features

| Feature         | XML Element  | Description   | Documentation  |
|-----------------|--|---|--|
| amp             | <amp<br>xmlns='http://jabber.org/features/amp'>              | Support for Advanced Message Processing                     | <a href="#">XEP-0079: Advanced Message Processing</a>                |
| compress        | <compression<br>xmlns='http://jabber.org/features/compress'> | Support for Stream Compression                              | <a href="#">XEP-0138: Stream Compression</a>                         |
| iq-auth         | <auth xmlns='http://jabber.org/features/iq-auth'>            | Support for Non-SASL Authentication                         | <a href="#">XEP-0078: Non-SASL Authentication</a>                    |
| iq-register     | <register<br>xmlns='http://jabber.org/features/iq-register'> | Support for In-Band Registration                            | <a href="#">XEP-0077: In-Band Registration</a>                       |
| bind            | <bind xmlns='urn:ietf:params:xml:ns:xmpp-bind'>              | Support for Resource Binding                                | <a href="#">RFC 6120: XMPP Core</a>                                  |
| mechanisms      | <mechanisms<br>xmlns='urn:ietf:params:xml:ns:xmpp-sasl'>     | Support for Simple Authentication and Security Layer (SASL) | <a href="#">RFC 6120: XMPP Core</a>                                  |
| session         | <session<br>xmlns='urn:ietf:params:xml:ns:xmpp-session'>     | Support for IM Session Establishment                        | <a href="#">RFC 6121: XMPP IM</a>                                    |
| starttls        | <starttls<br>xmlns='urn:ietf:params:xml:ns:xmpp-tls'>        | Support for Transport Layer Security (TLS)                  | <a href="#">RFC 6120: XMPP Core</a>                                  |
| bidi            | <bidi xmlns='urn:xmpp:bidi'>                                 | Support for Bidirectional Server-to-Server Connections      | <a href="#">XEP-0288: Bidirectional Server-to-Server Connections</a> |
| Server Dialback | <dialback<br>xmlns='urn:xmpp:features:dialback'>             | Support for Server Dialback with dialback errors            | <a href="#">XEP-0220: Server Dialback</a>                            |
| sm              | <sm xmlns='urn:xmpp:sm:3'>                                   | Support for Stream Management                               | <a href="#">XEP-0198: Stream Management</a>                          |

# Security Stuff

- Start-TLS
  - Prove the identity of the server
  - Prove the identity of the user (optional)
  - Encryption
  - Data integrity
- SASL(Simple Authentication and Security Layer protocol) ([RFC 4422](#))
  - Authentication
  - Optional encryption (rarely used)
  - Pluggable (e.g. passwords, Kerberos, X.509, SAML, etc.)

# Stanzas

- All have `to='JID'` and `from='JID'` addresses
  - To gives destination
  - From added by local server
- Each stanza routed separately
- All contents of stanza passed along
- Extend with any XML from your namespace
- Different types for delivery semantics

`<message/>`: one direction, one recipient

`<presence/>`: one direction, publish to many

`<iq/>`: "info/query", request/response

See details  
next page

# Message

- Example:

```
<message xml:lang='en'  
        to='romeo@example.net'  
        from='juliet@example.com/balcony'  
        type='chat'>  
  <body>Wherefore art thou, Romeo?</body>  
</message>
```

- Types: chat, groupchat, headline, error
- Body: plain text
- XHTML IM: [XEP-0071](#)

# Presence express an entity's current network availability

- Example:

```
<presence>  
  <show>dnd</show>  
  <status>Meeting</status>  
  <priority>1</priority>  
</presence>
```

- Show: chat, available, away, xa, dnd
- Status: Human-readable text
- Priority: Which resource "most available"?

# IQ Request

a structured request-response mechanism

- Example:

```
<iq type='get'
    id='roster_1'>
  <query xmlns='jabber:iq:roster' />
</iq>
```

- Type: get, set, result, error
- ID: track the corresponding response
- Query/Namespace: what type of request?

# IQ Response (Roster)

- Example:

```
<iq type='result'
  id='roster_1'>
  <query xmlns='jabber:iq:roster'>
    <item jid='romeo@example.net'
      name='Romeo'
      subscription='both'>
      <group>Friends</group>
    </item>
  </query>
</iq>
```

- Type: response
- ID matches request
- Subscription state: none, to, from, both



# Subscribing to Presence

- Send a subscription request:

```
<presence to='juliet@example.com'  
  type='subscribe' />
```

- Approving a request:

```
<presence to='romeo@example.net'  
  type='subscribed' />
```

- Every time you change a subscription, you get a "roster push":

```
<iq type='set'  
  <query xmlns='jabber:iq:roster'  
    <item jid='romeo@example.net'  
      subscription='from' />  
  </query>  
</iq>
```

# Extensibility Example: Message

- Use a new namespace
- Key: if you don't understand it, ignore it
- Example, CAP, [XEP-0127](#):

## Common Alerting Protocol (CAP) Over XMPP

```
<message to='weatherbot@jabber.org'
        from='KSTO@NWS.NOAA.GOV'>
  <alert xmlns='http://www.incident.com/cap/1.0'>
    <identifier>KSTO1055887203</identifier>
    <sent>2003-06-17T14:57:00-07:00</sent>
    <info>
      <category>Met</category>
      <event>SEVERE THUNDERSTORM</event>
    ...
  </info>
</alert>
</message>
```

# Extensibility Example: Presence

- Keep presence stanzas small
- Example: Entity Capabilities, [XEP-0115](#):  

```
<presence from='bard@shakespeare.lit/globe'>  
  <c xmlns='http://jabber.org/protocol/caps'  
    hash='sha-1'  
    node='http://www.chatopus.com'  
    ver='zHyEOgxTrkpSdGcQKH8EFPLsriY=' />  
</presence>
```

- Ver attribute is hash of all features of this client
- Hash -> Feature list is cached

It defines an XMPP protocol extension for broadcasting and dynamically discovering client, device, or generic entity capabilities.

# XMPP Extensions

- Many already exist: <http://www.xmpp.org/extensions/>
- Add new ones
  - Custom: use a namespace you control, make up protocol
  - Standardized: write a XEP.

| Number                   | Name                           | Type            | Status | Date       |
|--------------------------|--------------------------------|-----------------|--------|------------|
| <a href="#">XEP-0001</a> | XMPP Extension Protocols       | Procedural      | Active | 2016-11-16 |
| <a href="#">XEP-0002</a> | Special Interest Groups (SIGs) | Procedural      | Active | 2002-01-11 |
| <a href="#">XEP-0004</a> | Data Forms                     | Standards Track | Final  | 2007-08-13 |
| <a href="#">XEP-0009</a> | Jabber-RPC                     | Standards Track | Final  | 2011-11-10 |
| <a href="#">XEP-0012</a> | Last Activity                  | Standards Track | Final  | 2008-11-26 |
| <a href="#">XEP-0019</a> | Streamlining the SIGs          | Procedural      | Active | 2002-03-20 |
| <a href="#">XEP-0030</a> | Service Discovery              | Standards Track | Final  | 2017-10-03 |
| <a href="#">XEP-0047</a> | In-Band Bytestreams            | Standards Track | Final  | 2012-06-22 |
| <a href="#">XEP-0049</a> | Private XML Storage            | Historical      | Active | 2004-03-01 |
| <a href="#">XEP-0053</a> | XMPP Registrar Function        | Procedural      | Active | 2016-12-01 |
| <a href="#">XEP-0054</a> | vcard-temp                     | Historical      | Active | 2008-07-16 |
| <a href="#">XEP-0055</a> | Jabber Search                  | Historical      | Active | 2009-09-15 |

# Federation: DNS

- Starts with: non-local domain in to address

```
# _service._proto.name.  TTL   class SRV priority weight port target.  
_sip._tcp.example.com.  86400 IN     SRV 10      60    5060 bigbox.example.com.  
_sip._tcp.example.com.  86400 IN     SRV 10      20    5060 smallbox1.example.com.  
_sip._tcp.example.com.  86400 IN     SRV 10      20    5060 smallbox2.example.com.  
_sip._tcp.example.com.  86400 IN     SRV 20      0     5060 backupbox.example.com.
```

- Look up this DNS SRV record: (service record)  
`_xmpp-server._tcp.domain`

- Example: jabber.com:  
`10 0 5269 jabber.com.`

- Priority: Which one to try first if multiple
- Weight: Within a priority, what percentage chance?
- Port: TCP port number
- Target: Machine to connect to

# Federation: Security

- Old-style: dialback
  - Connect back to domain claimed by initiator
  - Check secret claimed by initiator
  - "Someone said they were example.com; was that you?"
- New-style: Mutual TLS
  - Initiator presents "client" certificate
  - Responder presents "server" certificate
  - Both certificates signed by trusted CA
- All stanzas *must* have from with correct domain

# Bandwidth minimization

- TLS compression
  - Not implemented in all SSL/TLS stacks
  - Some want compression w/o encryption
- XEP-0138: Stream Compression
  - Defines zlib mechanism (2-3x or more compression)
  - Others can be added
  - Concern: battery drain vs. radio transmission
- XEP-0198: Stanza Acknowledgements
  - Quick reconnects
  - Avoid re-synchronizing state on startup
- Partial rosters
- Privacy lists
- Others being pursued

# Latency

- Most critical on startup
  - Several handshakes and stream restarts
  - Can be minimized by client assuming server configuration
  - Example: don't wait for `<stream:features>`
- Once running
  - Stanza size matters: try to stay under 8kB, take larger blocks out of band if possible
  - Configure federation to keep links open, first stanza will be slow
  - Beware of DoS protection, "karma"



# Reading List

## ■ RFCs

- [6120](#): Core
- [6121](#): IM & Presence
- [5122](#): XMPP URIs

## ■ XEP highlights

- [4](#): Forms
- [30](#): Disco
- [45](#): Chat rooms
- [60](#): Pub/Sub
- [71](#): XHTML
- [115](#): Capabilities
- [163](#): PEP

# Advantages of XMPP

- The primary advantage is XMPP's **decentralized** nature.
- XMPP works similar to email, operating across a **distributed network** of transfer agents **rather than relying on a single, central server or broker** (as CoAP and MQTT do).
- As with email, it's easy for anyone to run their own XMPP server, allowing device manufacturers and API operators to create and manage their own network of devices.
- And because anyone can run their own server, if security is required, that server could be isolated on a company intranet behind secure authentication protocols using built-in TLS encryption.

# Disadvantages of XMPP

- One of the largest flaws is the **lack of end-to-end encryption**. While there are many use cases in which encryption may not yet be necessary, most IoT devices will ultimately need it. The lack of end-to-end encryption is a major downside for IoT manufacturers.

[https://wiki.xmpp.org/web/XMPP\\_E2E\\_Security](https://wiki.xmpp.org/web/XMPP_E2E_Security)

- Another downside is the **lack of Quality of Service (QoS)**. Making sure that messages are delivered is even more important in the IoT world than it was in the instant messaging world. If your alarm system doesn't receive the message to turn itself on, then that vacation you've been planning could easily be ruined.